



# Load Balancer

## Using P4 and BMV2 hash extern

(<https://github.com/Tigohsr/loadBalancer-P4>)


Felipe Novais

Thiago Henrique S Rodrigues

# topo.txt

```
switches 4  
hosts 4  
h1 s1  
s1 s2  
s1 s3  
s1 s4  
s2 h2  
s2 s4  
s3 h3  
s3 s4  
s4 h4
```

```
topo.txt
```



# topo.py

```
105 sw_addr = ["10.0.%d.254" % n for n in xrange(nb_hosts)]
```

```
59     for h in xrange(nb_hosts):
60         host = self.addHost('h%d' % (h + 1),
61                             ip = "10.0.%d.%d/24" % (h+1, h+1),
62                             mac = '00:04:00:00:00:%02x' % h)
63
64     i = 0
```

```
178     if (i == 0):
179         print "Running lb_s1.txt"
180         cmd = [args.cli, "--json", args.json,
181               "--thrift-port", str(_THRIFT_BASE_PORT + i)]
182         with open("lb_s1.txt", "r") as f:
183             print " ".join(cmd)
184             try:
185                 output = subprocess.check_output(cmd, stdin = f)
186                 #print ("Debugging switch match-action table")
187                 print output
188             except subprocess.CalledProcessError as e:
189                 print e
190                 print e.output
191         print "*****"
```


# Tables in Control Plane



table_set_default ipv4_lpm drop		
table_set_default lb_hash_exact drop		
~		
~		
~		
command_s1.txt	1,1	All
table_set_default ipv4_lpm drop		
table_add ipv4_lpm ipv4_forward 10.0.2.2/32 => 00:04:00:00:00:01 2		
~		
~		
command_s2.txt	1,1	All
table_set_default ipv4_lpm drop		
table_add ipv4_lpm ipv4_forward 10.0.3.3/32 => 00:04:00:00:00:02 2		
~		
~		
command_s3.txt	1,1	All
table_set_default ipv4_lpm drop		
table_add ipv4_lpm ipv4_forward 10.0.4.4/32 => 00:04:00:00:00:03 4		
~		
~		
command_s4.txt	1,1	All

# ports.py 1/2

```
1 import os, binascii, sys
2
3 # File that will store control plane table entries
4 s1_rules = open('lb_s1.txt', 'w')
5
6 # IPs in hex
7 h1 = '0a000101' # 10.0.1.1
8 s1 = '0a0001fe' # 10.0.1.254
9
10 h2 = '0a000202' # 10.0.2.2
11 h3 = '0a000303' # 10.0.3.3
12 h4 = '0a000404' # 10.0.4.4
13
14 # Original locations for destination from h1 to s1
15 orig_destinations = [h2, h3, h4]
16
17 # Original egress ports from h1 to s1
18 s1_ports = [2, 3, 4]
19
20 # TCP ports from 1000 to 1499
21 dports = range(1000, 1500)
22
23 # Fixed source port to 1234
24 sport = '04d2' # 1234
25
26 tcp_proto = '06' # IP Protocol
27
28 ports = []
29 rules = []
```



# ports.py 2/2

```
31 print(" **** Generating rules for load balancer on S1... **** ")
32 for dport in dports:
33     # Initialize copy of original destination and ports on s1
34     if len(ports) == 0:
35         ports = s1_ports.copy()
36         destinations = orig_destinations.copy()
37
38     # Cryptographically secure random function to select a random port on s1
39     port = int(binascii.hexlify(os.urandom(2 ** 16)), 16) % len(ports)
40
41     # Pop port from list to provide a certain fairness during port selection
42     s1_port = str(ports.pop(port))
43     destination = destinations.pop(port)
44
45     # TCP-tuple that will be hashed
46     tcp_tuple = h1 + s1 + sport + format(dport, '0>4x') + tcp_proto
47
48     # CRC32 hashing the TCP-tuple
49     crc32_hash = str(binascii.crc32(binascii.a2b_hex(tcp_tuple)))
50
51     # Append to a list the rule that will be an entry on the control plane
52     rules.append('table_add lb_hash_exact lb_hash_forward ' + crc32_hash + '
=> 0x' + destination + ' ' + s1_port)
53
54
55 # Writing list to file
56 s1_rules.write('\n'.join(rule for rule in rules))
57
58 print(" **** Generated rules for load balancer on S1 **** ")
```

Iteration #1:

[ 2, 3, 4]  
[ h2, h3, h4]

Port - randomized index: 2

Iteration #2:

[ 2, 4]  
[ h2, h4]

Port - randomized index: 1

Iteration #3:

[ 4]  
[h4]

Port - randomized index: 1

Iteration #4:

[ 2, 3, 4]  
[ h2, h3, h4]

...



# run\_task.sh



```
THIS_DIR=$( cd "$( dirname "${BASH_SOURCE[0]}" )" && pwd )

source $THIS_DIR/../env.sh

P4C_BM_SCRIPT=/usr/local/bin/p4c/build/p4c

SWITCH_PATH=/usr/local/bin/simple_switch

CLI_PATH=$BMV2_PATH/tools/runtime_CLI.py
#
$P4C_BM_SCRIPT --target bmv2-ss-p4org -x p4-16 p4src/assignment.p4 -o p4prog
# This gives libtool the opportunity to "warm-up"
sudo $SWITCH_PATH >/dev/null 2>&1
sudo PYTHONPATH=$PYTHONPATH:$BMV2_PATH/mininet/ python3 ports.py
sudo PYTHONPATH=$PYTHONPATH:$BMV2_PATH/mininet/ python topo.py \
    --behavioral-exe $SWITCH_PATH \
    --json p4prog/assignment.json \
    --cli $CLI_PATH
```

# Generated Control Plane (lb\_s1.txt)



table_add	lb_hash_exact	lb_hash_forward	451078402 =>	0x0a000202	2
table_add	lb_hash_exact	lb_hash_forward	835697345 =>	0x0a000404	4
table_add	lb_hash_exact	lb_hash_forward	685017984 =>	0x0a000404	4
table_add	lb_hash_exact	lb_hash_forward	1737825607 =>	0x0a000303	3
table_add	lb_hash_exact	lb_hash_forward	2123246598 =>	0x0a000202	2
table_add	lb_hash_exact	lb_hash_forward	1436778437 =>	0x0a000202	2
table_add	lb_hash_exact	lb_hash_forward	1287147140 =>	0x0a000303	3
table_add	lb_hash_exact	lb_hash_forward	2179088410 =>	0x0a000404	4
table_add	lb_hash_exact	lb_hash_forward	2566484315 =>	0x0a000303	3
table_add	lb_hash_exact	lb_hash_forward	3017024152 =>	0x0a000404	4
table_add	lb_hash_exact	lb_hash_forward	2865697753 =>	0x0a000202	2
table_add	lb_hash_exact	lb_hash_forward	3851322654 =>	0x0a000303	3
table_add	lb_hash_exact	lb_hash_forward	4237669471 =>	0x0a000404	4
table_add	lb_hash_exact	lb_hash_forward	3619219356 =>	0x0a000202	2
table_add	lb_hash_exact	lb_hash_forward	3466843869 =>	0x0a000404	4
table_add	lb_hash_exact	lb_hash_forward	1228655122 =>	0x0a000202	2
table_add	lb_hash_exact	lb_hash_forward	1344338771 =>	0x0a000303	3
table_add	lb_hash_exact	lb_hash_forward	2064492688 =>	0x0a000202	2
table_add	lb_hash_exact	lb_hash_forward	1645648337 =>	0x0a000404	4
table_add	lb_hash_exact	lb_hash_forward	760680214 =>	0x0a000303	3
table_add	lb_hash_exact	lb_hash_forward	877411927 =>	0x0a000303	3
table_add	lb_hash_exact	lb_hash_forward	526478740 =>	0x0a000202	2
table_add	lb_hash_exact	lb_hash_forward	108682453 =>	0x0a000404	4
table_add	lb_hash_exact	lb_hash_forward	2417080098 =>	0x0a000303	3
table_add	lb_hash_exact	lb_hash_forward	2299168355 =>	0x0a000202	2
table_add	lb_hash_exact	lb_hash_forward	2720519584 =>	0x0a000404	4
table_add	lb_hash_exact	lb_hash_forward	3141330145 =>	0x0a000303	3
table_add	lb_hash_exact	lb_hash_forward	4101861926 =>	0x0a000202	2
table_add	lb_hash_exact	lb_hash_forward	3982902119 =>	0x0a000404	4
table_add	lb_hash_exact	lb_hash_forward	3326808228 =>	0x0a000404	4
table_add	lb_hash_exact	lb_hash_forward	3746570725 =>	0x0a000303	3



# Defines and types



```
6 const bit<16> TYPE_IPV4 = 0x800;
7
8 // Max integer for a 32bits variable
9 #define MAX_32BITS 4294967295
10
11 // IP protocol field that represents TCP
12 #define TCP_PROTO 6
13
14 // Device IP in hex
15 #define S1 0x0a0001fe
16
```

```
49 header tcp_t {
50     bit<16> srcPort;
51     bit<16> dstPort;
52     bit<32> seqNo;
53     bit<32> ackNo;
54     bit<4> dataOffset;
55     bit<3> res;
56     bit<3> ecn;
57     bit<6> ctrl;
58     bit<16> window;
59     bit<16> checksum;
60     bit<16> urgentPtr;
61 }
62
63 struct metadata {
64     bit<32> hashedTuple;
65 }
```

# Parsers



```
87     state parse_ethernet {
88         packet.extract(hdr.ethernet);
89         transition select(hdr.ethernet.etherType) {
90             TYPE_IPV4: parse_ipv4;
91             default: accept;
92         }
93     }
94
95
96     state parse_ipv4 {
97         packet.extract(hdr.ipv4);
98         transition select(hdr.ipv4.protocol) {
99             TCP_PROTO: parse_tcp;
100             default: accept;
101         }
102     }
103
104     state parse_tcp {
105         packet.extract(hdr.tcp);
106         transition accept;
107     }
```

# Tables in P4

```
145     table ipv4_lpm {
146         key = {
147             hdr.ipv4.dstAddr: lpm;
148         }
149         actions = {
150             ipv4_forward;
151             drop;
152             NoAction;
153         }
154         size = 1024;
155         default_action = NoAction();
156     }
157
158     table lb_hash_exact {
159         key = {
160             meta.hashTuple: exact;
161         }
162         actions = {
163             lb_hash_forward;
164             drop;
165             NoAction;
166         }
167         size = 1024;
168         default_action = NoAction();
169     }
```

# Apply

```
171     apply {
172         if (hdr.ipv4.isValid()) {
173             // Checks if dstAddr is S1
174             if (hdr.ipv4.dstAddr == S1) {
175                 // Hashing TCP-tuple using CRC32
176                 hash(meta.hashedTuple,
177                     HashAlgorithm.crc32,
178                     (bit<32>)0,
179                     {hdr.ipv4.srcAddr, hdr.ipv4.dstAddr, hdr.tcp.srcPort, hdr.tcp.dstPort, hdr.ipv4.protocol},
180                     (bit<32>)MAX_32BITS);
181                 lb_hash_exact.apply();
182             } else {
183                 ipv4_lpm.apply();
184             }
185         } else {
186             drop();
187         }
188     }
189 }
```

output  
hash algorithm  
minimum value  
data  
maximum value

```
// Hashing TCP-tuple using CRC32
hash(meta.hashedTuple,
    HashAlgorithm.crc32,
    (bit<32>)0,
    {hdr.ipv4.srcAddr, hdr.ipv4.dstAddr, hdr.tcp.srcPort, hdr.tcp.dstPort, hdr.ipv4.protocol},
    (bit<32>)MAX_32BITS);
```

## Atividade 2 - Balanceador de Carga.pdf:

- um fluxo é caracterizado pelos 5 campos: IP origem, IP destino, porta origem, porta destino e protocolo;



# Actions



```
132  action lb_hash_forward(ip4Addr_t dstAddr, egressSpec_t port) {
133      standard_metadata.egress_spec = port;
134      hdr.ipv4.dstAddr = dstAddr;
135      hdr.ipv4.ttl = hdr.ipv4.ttl - 1;
136  }
137
138  action ipv4_forward(macAddr_t dstAddr, egressSpec_t port) {
139      standard_metadata.egress_spec = port;
140      hdr.ethernet.srcAddr = hdr.ethernet.dstAddr;
141      hdr.ethernet.dstAddr = dstAddr;
142      hdr.ipv4.ttl = hdr.ipv4.ttl - 1;
143  }
```

# send.py



```
27 def main():
28     
29     if len(sys.argv)<1:
30         print 'pass 1 arguments: <destination_port>'
31         exit(1)
32
33     s1 = "10.0.1.254"
34     addr = socket.gethostbyname(s1)
35     iface = get_if()
36     print "sending on interface %s to %s on dport: %s" % (iface, str(addr), sys.argv[1])
37
38     pkt = Ether(src=get_if_hwaddr(iface), dst='ff:ff:ff:ff:ff:ff');
39
40     pkt = pkt / IP(src=get_if_addr(iface), dst=addr) / TCP(dport=int(sys.argv[1]), sport=1234)
41
42     pkt.show2()
43     sendp(pkt, iface=iface, verbose=False)
44
45
46 if __name__ == '__main__':
47     main()
```



# DEMO



# Thank you!